

SRAVAN POTNURU

AI / ML Engineer

PROFILE SUMMARY

AI/ML Engineer with 2 years of experience building production-grade video intelligence and GenAI applications. Experienced in FastAPI microservices, Triton Inference Server, AWS cloud services, computer vision, and LLM-powered systems. Currently involved in LLM fine-tuning and post-training R&D, while contributing to a US client Enterprise AI Harness for intelligent interaction with SAP data and enterprise documents. More technical version

WORK EXPERIENCE

AI/ML Engineer

Parabola9

12/2024 – Present | Mangalagiri

- Designed and deployed a production video analytics platform processing 10+ live CCTV streams through FastAPI microservices, enabling real-time AI inference and event detection.
- Built a serverless video ingestion pipeline using AWS Lambda, Kinesis Video Streams and S3 to automatically ingest, process and store live CCTV footage, eliminating manual video handling.
- Optimized AI inference using NVIDIA Triton Inference Server, reducing model serving latency and enabling real-time object detection across multiple concurrent streams.
- Developed a computer vision pipeline combining YOLO, BoT-SORT and InsightFace to perform real-time person detection, tracking and face recognition for surveillance applications.
- Containerized and deployed 5+ AI microservices on AWS EC2 using Docker, reducing deployment time and improving release consistency.
- Built an AI-powered warehouse surveillance platform capable of detecting PPE compliance and security events in real time, improving operational visibility for government facilities.
- Collaborate directly with CEO & CTO (US stakeholders) for development plan, performance improvements, and deployment decisions
- Deployed GPU-enabled AI inference services on Akash Network (Linux/AMD64), reducing infrastructure cost while supporting production-scale model serving.
- Built a real-time video ingestion pipeline using Apache NiFi to ingest live streams, automate video chunking and eliminate manual ingestion processes.

CLIENT PROJECT: RICH'S

02/2026 – Present | Mangalagiri

- Enhanced enterprise knowledge graph pipelines powering AI search across internal documentation, improving retrieval quality and contextual responses.
- Developed FastAPI APIs supporting enterprise document ingestion, semantic search and AI-driven knowledge retrieval for internal business users.
- Built Azure Data Factory ETL pipelines ingesting Azure DevOps work items and Azure File Share documents into Bronze, Silver and Gold layers, expanding enterprise AI knowledge coverage across multiple data sources.
- Executed large-scale Locust load testing to identify performance bottlenecks, improving backend scalability before production deployments.

PROJECTS

Project: ParaBot (AI Chat + File Summarizer) | FastAPI, React, OCR, Gemini | [GitHub](#)

Implemented an AI assistant that processes text, images, PDFs, and website content with summarization and Q&A capabilities.

Project: Stress Prediction Based on Sleeping Patterns | Python, ML, Streamlit | [GitHub](#)

Developed a machine learning model that predicts stress levels using sleep data and provides wellness suggestions.

Project: ScreenSaga (OTT Platform) | React JS, Spring RestAPI, MySQL | [GitHub](#) (frontend) | [GitHub](#) (Backend)

Built a full-stack OTT platform with movies, web series, trailers, personalized watchlists, secure authentication (email verification, password reset, account deletion), season/episode selection, Continue Watching, and a feedback system. Backend implemented using Spring Boot REST APIs, MySQL, JSON APIs, and SMTP for email confirmations.

Project: LogIQ - AI-Powered Project Monitor | FastAPI, React, PostgreSQL, Multi-LLM

(OpenAI/Anthropic/Azure/Ollama), Cloud Integration | [LogIQ Application](#)

Built a multi-tenant AI observability platform integrated into applications via a pip-installable Python SDK published on PyPI. Implemented log ingestion, AI-driven root-cause analysis with automatic rule-based fallback, a debugging chat assistant that identifies the source error file, multi-channel alerting (Slack/Teams/Email), and cloud-based monitoring (AWS/Azure/GCP) across backend APIs and a React dashboard.

CO-CURRICULAR

- Achieved NCC '[B](#)' Certificate and '[C](#)' Certificate.
- [Certificate of Merit for All India Essay Writing Event 2018.](#)
- [Certificate of Participation and Certificate of Appreciation for National Level Techno Management Fest](#)

CONTACT

- 📞 +91 9391872342
- ✉️ sravanpotnuru24@gmail.com
- 🌐 [Portfolio](#)
- 🌐 [LinkedIn](#)
- 🌐 [GitHub](#)

EDUCATION

B.Tech - Computer Science Engineering (CSE)

RGUKT IIIT Nuzvid

2021 - 2025

CGPA: 9.0

Pre- University Course (PUC) - Intermediate

RGUKT IIIT Nuzvid

2019 - 2021

CGPA: 9.62

Secondary School - SSC

Government High School, Srikakulam

2018 - 2019

CGPA: 10

TECH STACK

- **Languages:** Python, Java
- **Backend:** FastAPI, Spring Boot, REST APIs
- **AI/ML:** YOLO, InsightFace (Buffalo), RAG, LangGraph, Triton Inference Server
- **Cloud & DevOps:** AWS (EC2, S3, Lambda, Kinesis), Docker
- **Databases:** MySQL, Neo4j (Cypher Queries)
- **Data & Analytics:** Databricks, Azure Data Factory (ADF), Apache NiFi
- **Frontend:** React JS, HTML, CSS, Bootstrap
- **Testing & Performance:** Locust
- **Tools:** Git, Postman, Jupyter, VS Code